

# THE 23RD INTERNATIONAL SEMANTIC WEB CONFERENCE

November 11, 2024 – November 15, 2024

Live! Casino & Hotel Maryland

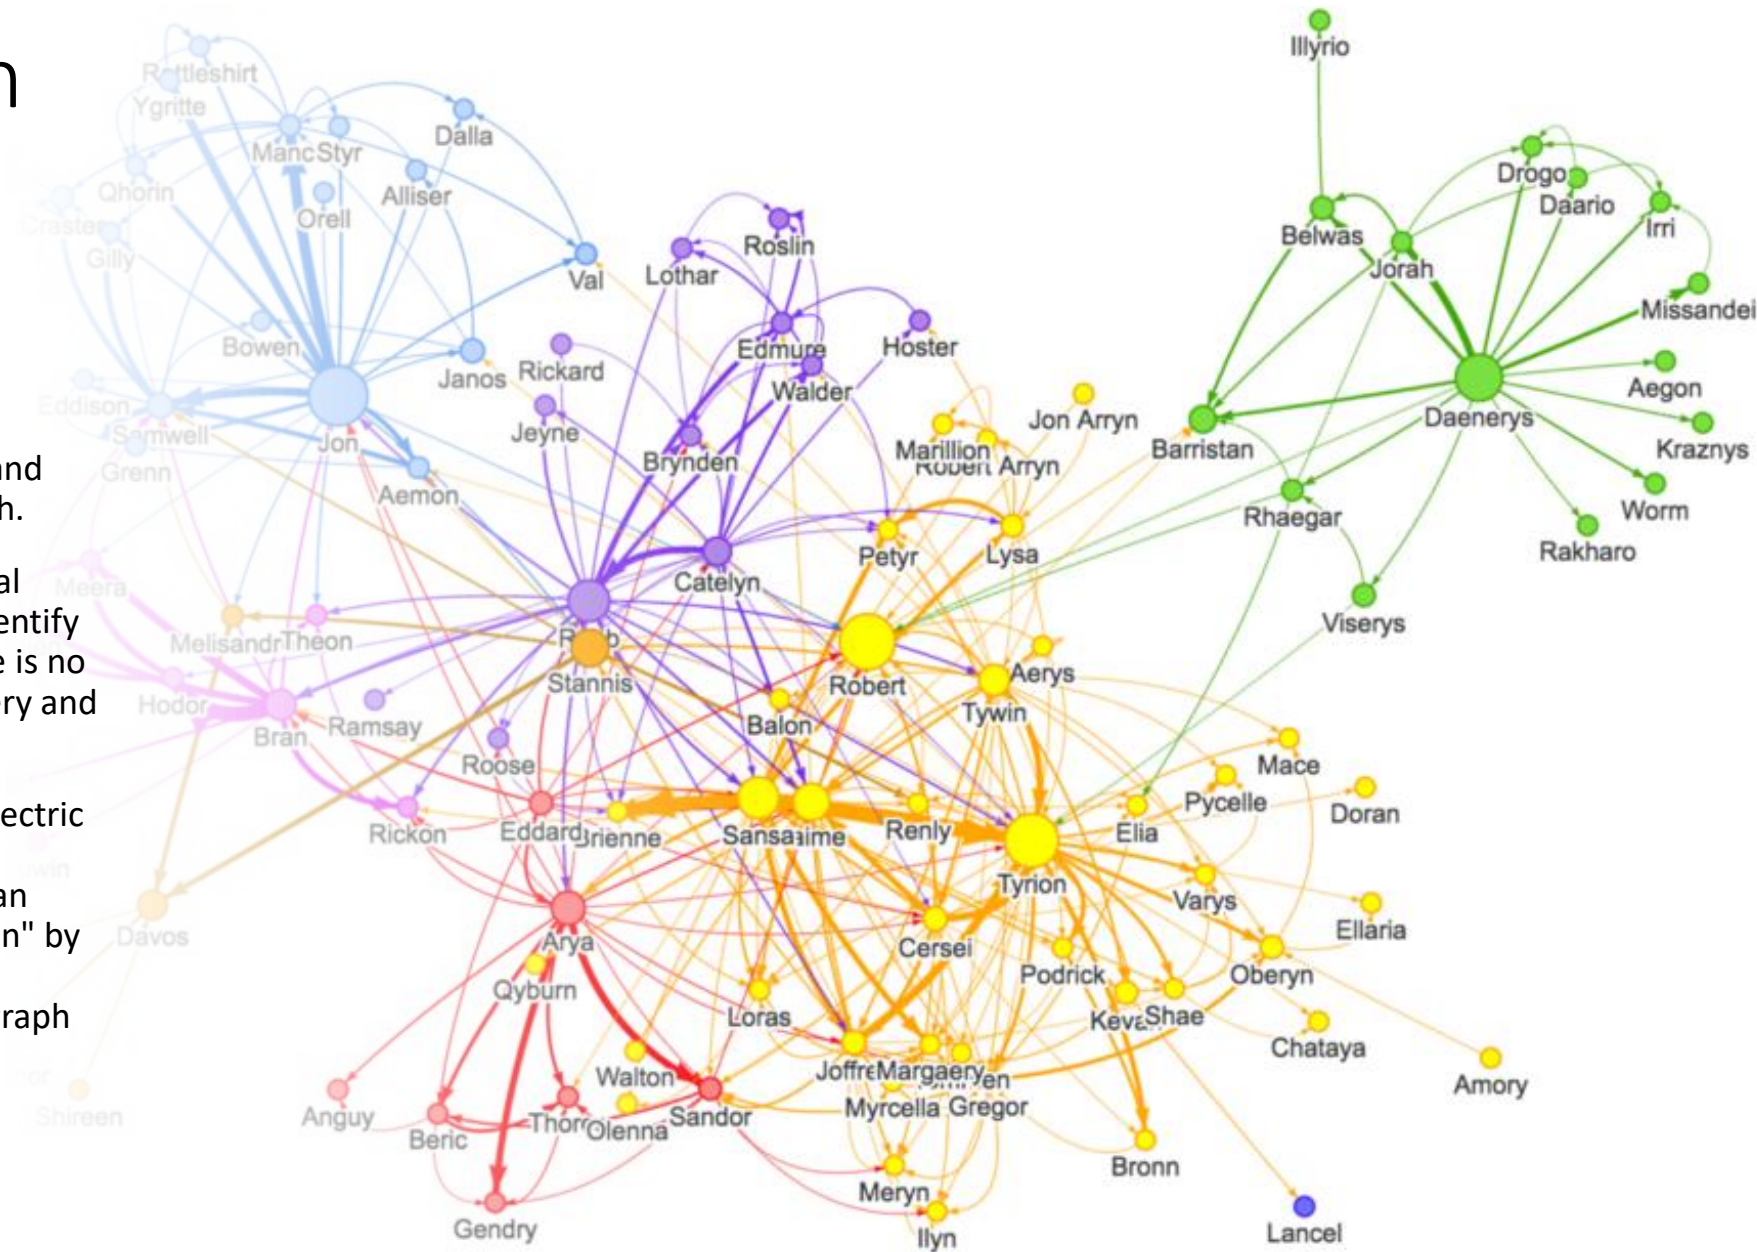
## Slot 3: Advanced Graph RAG Approaches

ISWC 2024

# Part 1: Introduction

# Graph RAG with Semantic Clustering

- This approach leverages clustering algorithms to group similar entities and concepts within the knowledge graph.
- This clustering enhances the retrieval process by allowing the system to identify related information even when there is no direct match between the user's query and the graph's entities.
- For instance, if a user asks about "electric vehicles," the system could retrieve information related to "Tesla," "Nissan Leaf," and "sustainable transportation" by recognizing their shared cluster membership within the knowledge graph




# GraphRAG

From Microsoft



# A Graph RAG Approach to Query-Focused Summarization

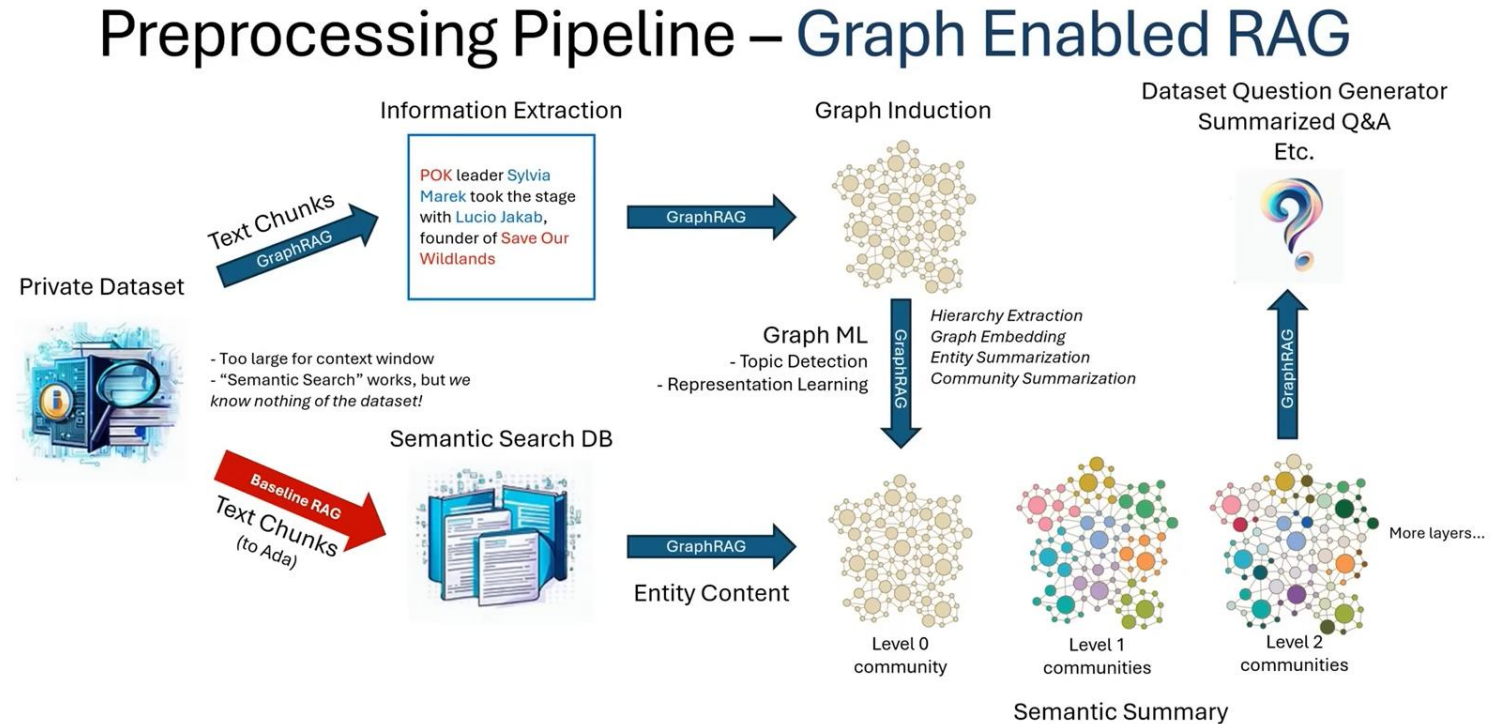
- Naive RAG Issue: Fails on global questions like “What are the main themes in the dataset?”
- Graph RAG Approach:
  - Extract knowledge graph from raw text
  - Build a community hierarchy
  - Generate community summaries
- Improved comprehensiveness and diversity of generated answers on large datasets



An LLM-generated knowledge graph built using GPT-4 Turbo, Microsoft, <https://microsoft.github.io/graphrag/>

# A Graph RAG Approach to Query-Focused Summarization

- Execute the indexing pipeline to extract and construct the knowledge graph.
- Convert artifacts into RDF triples, mapping them according to the ontology.
- Perform semantic searches within the RDF graph to retrieve detailed entity information.
- Extract relevant subgraphs for focused and efficient querying.



GraphRAG: LLM-Derived Knowledge Graphs for RAG, YouTube, uploaded by Alex Chao, May 4, 2024. Available at: <https://youtu.be/r09tJfON6kE>

# Part 2: G-Indexing

Creating a Knowledge Graph with LLMs for RAG systems

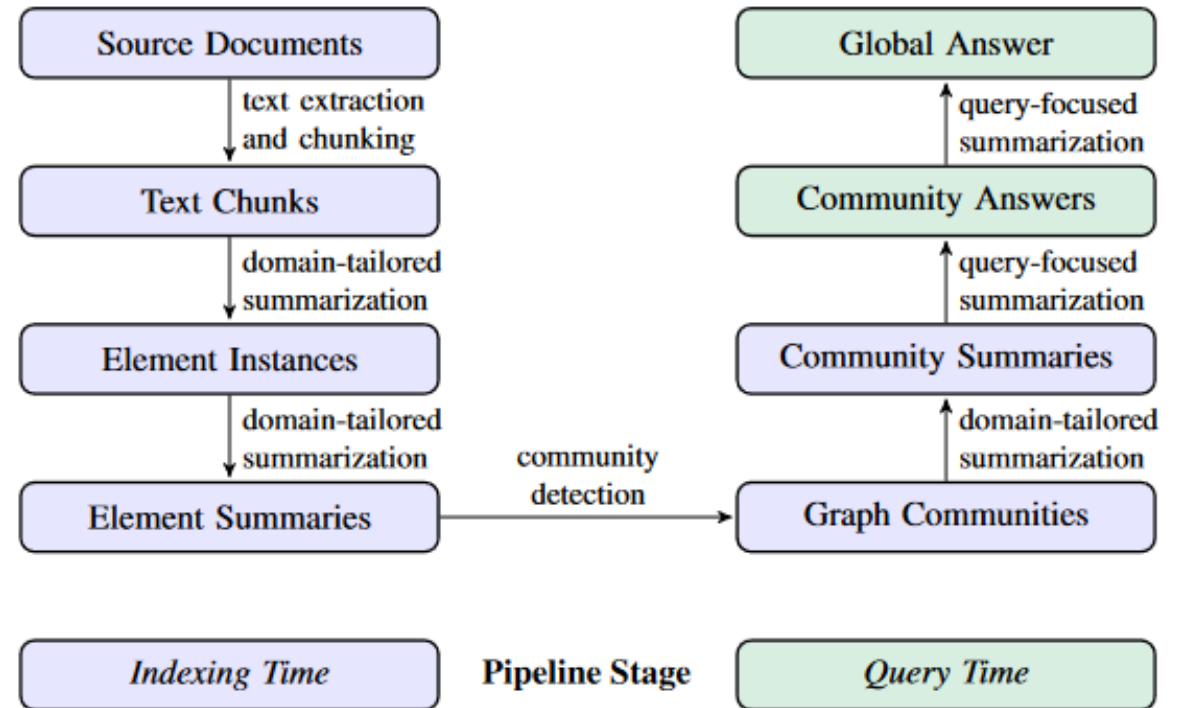
# Indexing Pipeline

## Knowledge Graph Construction:

- **LLM-Based:** Automatically identify entities and relationships.
- **Iterative Gleaning:** Multi-round processing ensures completeness.

## Community Summarization:

- Detect communities in the graph (e.g., using Leiden algorithm).
- Summarize each community for answering questions.

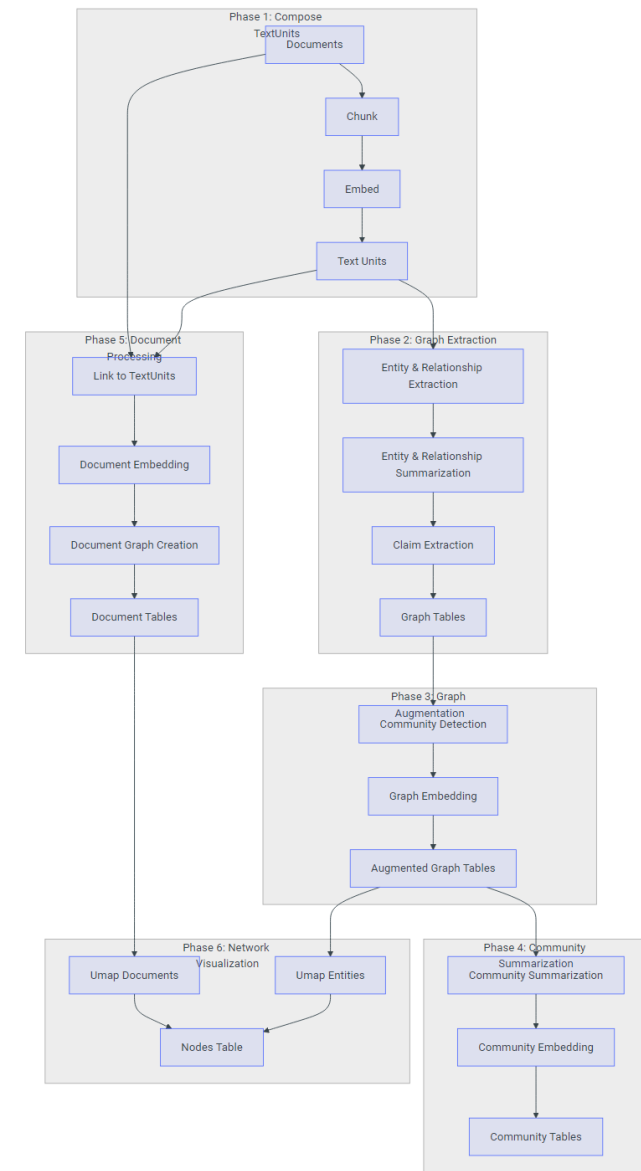
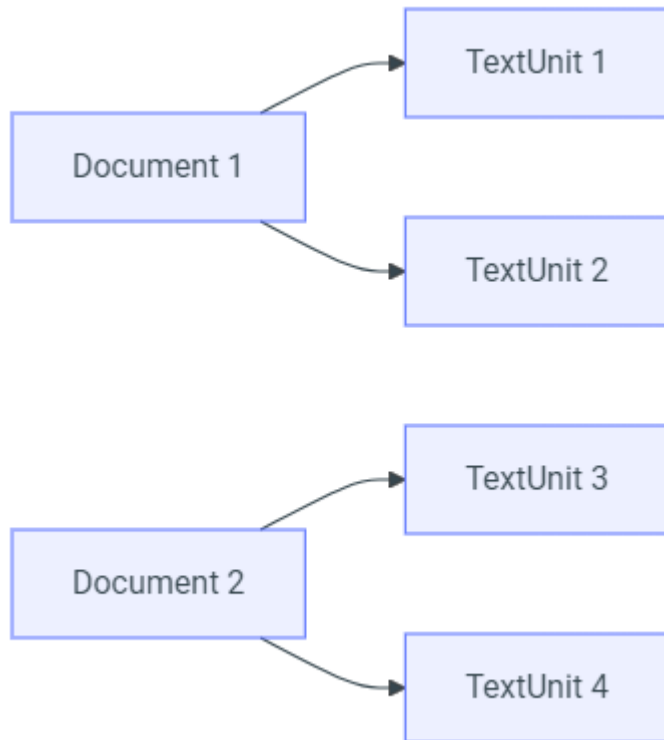


Edge, D., Trinh, H., Cheng, N., Bradley, J., Chao, A., Mody, A., Truitt, S., & Larson, J. (2024). From local to global: A graph RAG approach to query-focused summarization. arXiv. <https://arxiv.org/abs/2404.16130>



# Indexing Dataflow

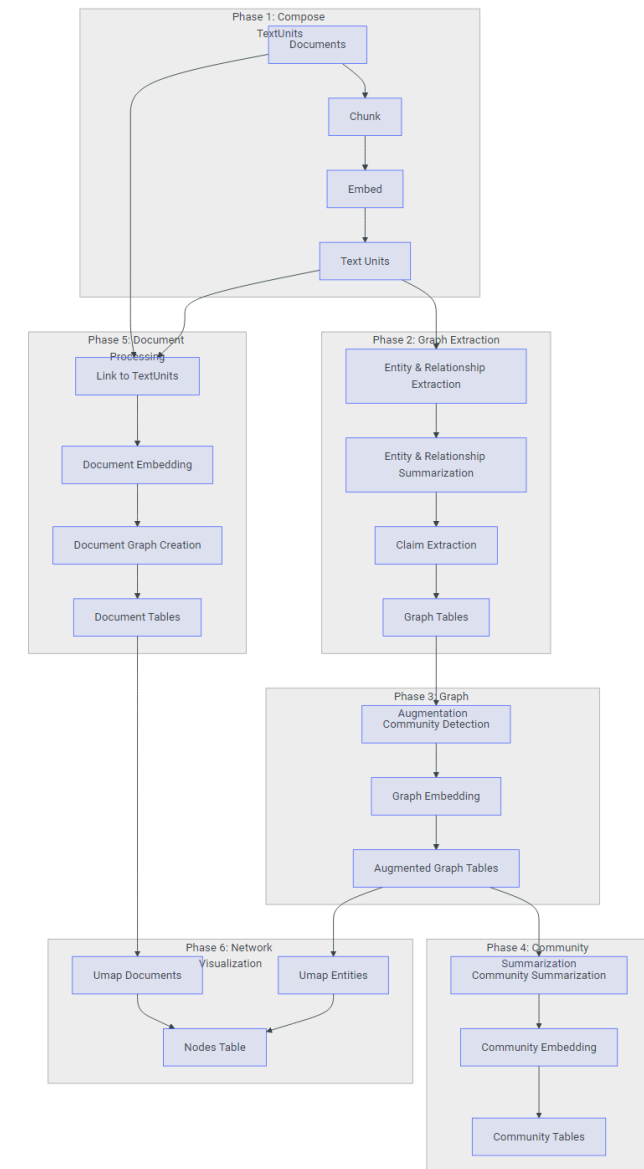
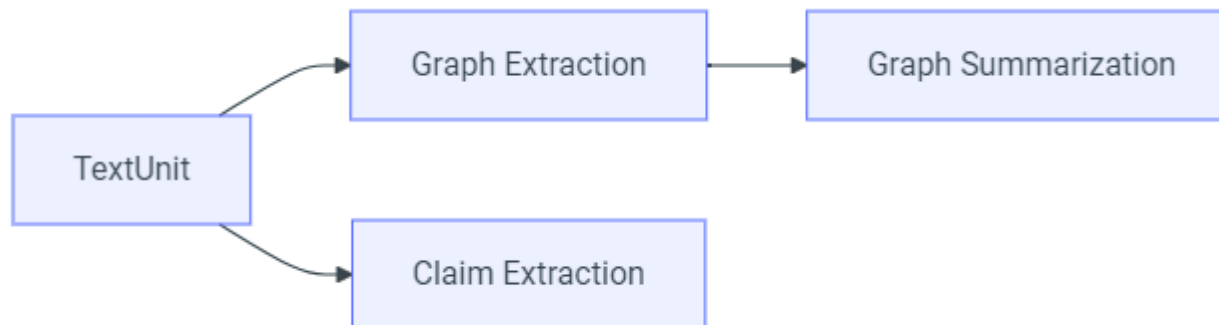
- Phase 1: Compose TextUnits
  - Transform input documents into *TextUnits*



The Default Configuration Workflow transforms text documents into the GraphRAG Knowledge Model, Microsoft, <https://microsoft.github.io/graphrag/>

# Indexing Dataflow

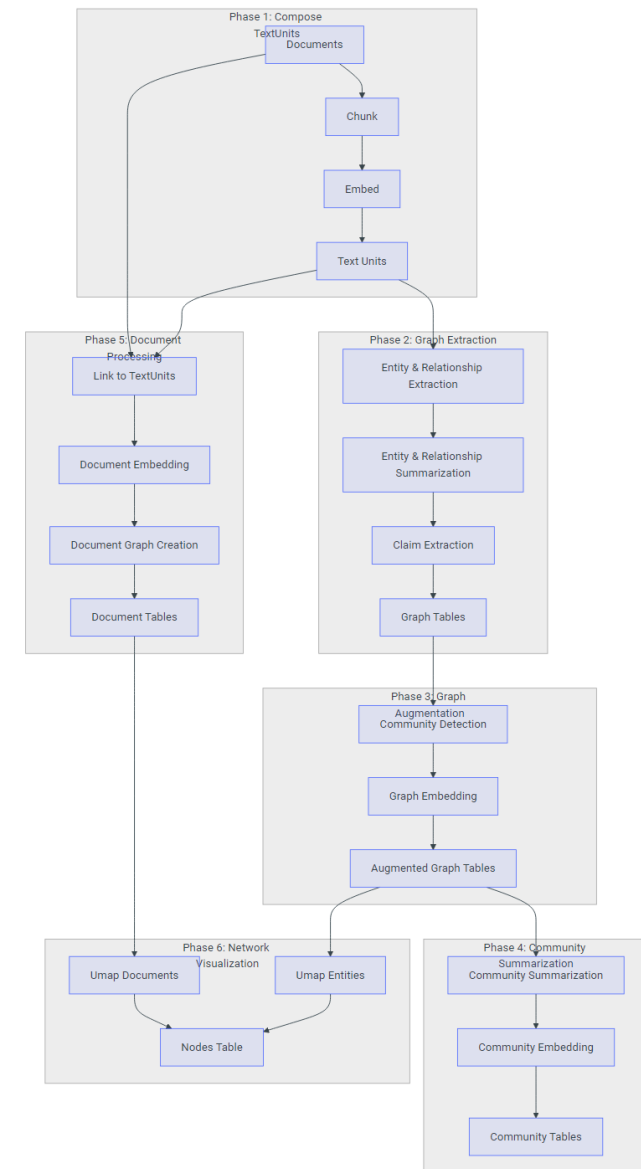

- Phase 2: Graph Extraction
  - Entity & Relationship Extraction
    - A list of entities with a name, type, and description
    - A list of relationships with a source, target, and description
  - Entity & Relationship Summarization
    - Summarize these lists into a single description per entity and relationship



The Default Configuration Workflow transforms text documents into the GraphRAG Knowledge Model, Microsoft, <https://microsoft.github.io/graphrag/>

# Indexing Dataflow


- Phase 3: Graph Augmentation
  - **Community Detection:** Generate a hierarchy of entity communities using the Hierarchical Leiden Algorithm
  - **Graph Embedding:** Generate a vector representation of our graph using the Node2Vec algorithm



The Default Configuration Workflow transforms text documents into the GraphRAG Knowledge Model, Microsoft, <https://microsoft.github.io/graphrag/>

# Indexing Dataflow


- Phase 4: Community Summarization
  - **Generate Community Reports:** Generate a summary of each community using the LLM
  - **Summarize Community Reports:** Each community report is then summarized via the LLM for shorthand use
  - **Community Embedding:** Generate a vector representation of communities



The Default Configuration Workflow transforms text documents into the GraphRAG Knowledge Model, Microsoft, <https://microsoft.github.io/graphrag/>

# Indexing Dataflow

- Phase 5: Document Processing
  - **Link to TextUnits:** Link each document to the text-units that were created in the first phase
  - **Document Embedding:** Generate a document embedding by averaging token-weighted, non-overlapping chunks to capture document relationships



The Default Configuration Workflow transforms text documents into the GraphRAG Knowledge Model, Microsoft, <https://microsoft.github.io/graphrag/>



# Running the Indexing pipeline

```
python -m graphrag.index --root ./ragtest
```

```
Loading csv files from ./input
loading 1 csv files
Total number of unfiltered csv rows: 4748
Final # of rows loaded: 4748
! Executing Pipeline...
├─ Loading Input (csv) - 1 files loaded (0 filtered) 100% 0:00:00 0:00:00
├─ Workflow: create_base_text_units
├─ Workflow: create_base_extracted_entities
└─ verb: entity_extract 99% 0:01:20 12:22:38
```

Indexing pipeline execution in GraphRAG, Microsoft, <https://microsoft.github.io/graphrag/>

# Knowledge Graph Visualization

- Artifacts from the Indexing Pipeline
  - The outputs of the indexing pipeline are a set of Parquet files, which serve as the knowledge base for the subsequent retrieval stage.
- For a quick overview of the graph structure, visit [GraphRAG-Visualizer](#).



The resulting knowledge graph visualization from GraphRAG, shown in the GraphRAG Visualizer.

# Part 3: RDF Adoptions

Adapting GraphRAG Artifacts to RDF

# RDF Adaptions

- We follow the steps below to transform GraphRAG artifacts into RDF.

**Data Ingestion:** Read Parquet files into Pandas Data Frames.



**RDF Graph Initialization:** Set up the RDF graph using rdflib.



**Ontology Definition:** Define classes and properties in RDF.



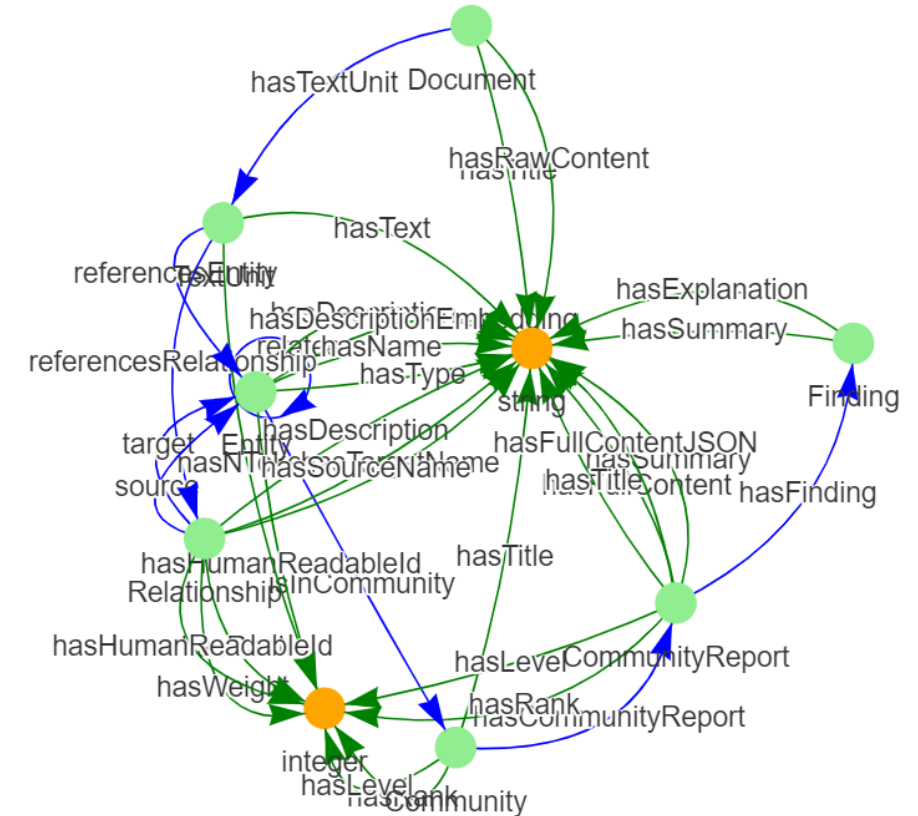
**Mapping Nodes and Relationships to RDF:** Convert each node and its relationships into RDF triples.



**Serializing the RDF Graph:** Export the RDF graph in turtle format.

# Ontology Visualization

Source	Relationship	Target
Entity	RELATES	Entity
Entity	IN_COMMUNITY	Community
Document	HAS_TEXTUNIT	TextUnit
Community	HAS_COMMUNITYREPORT	CommunityReport
CommunityReport	HAS_FINDING	Finding
TextUnit	REFERENCES_ENTITY	Entity



Visualization of the RDF Graph Ontology

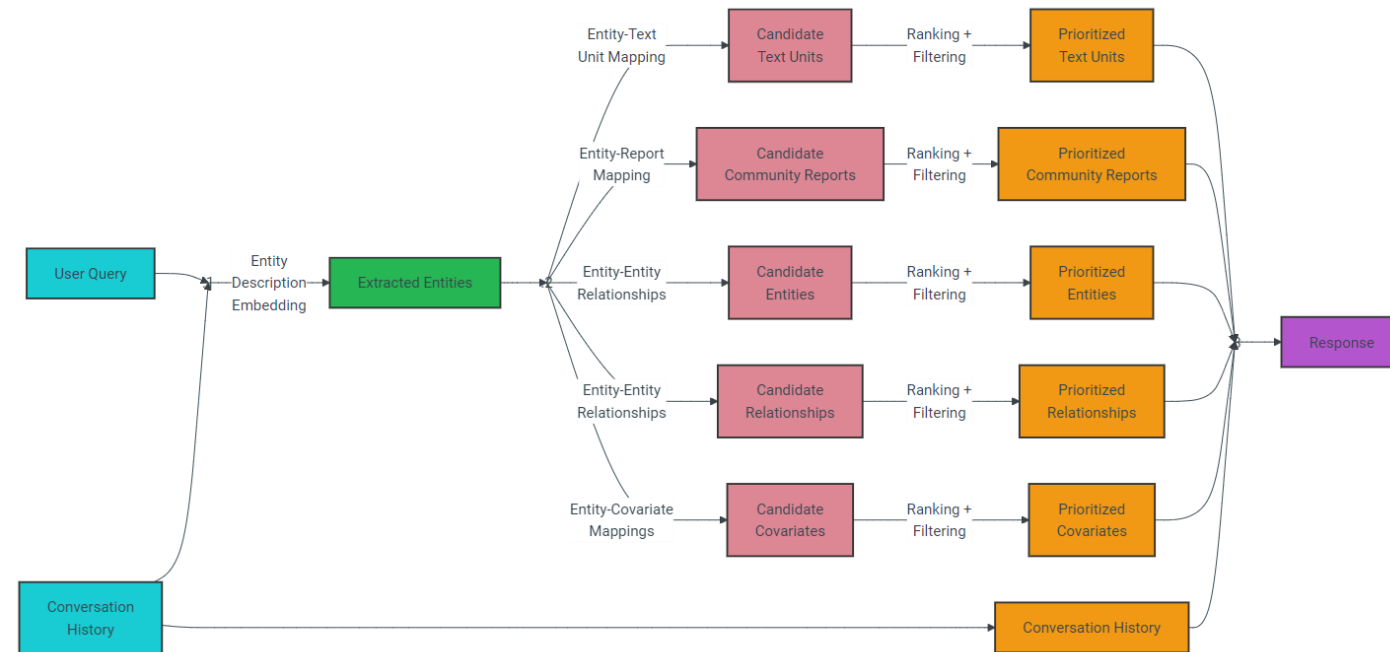


# Part 4: G-Retrieval & G-Generation

Using SPARQL with Knowledge Graphs for RAG

# Local Search

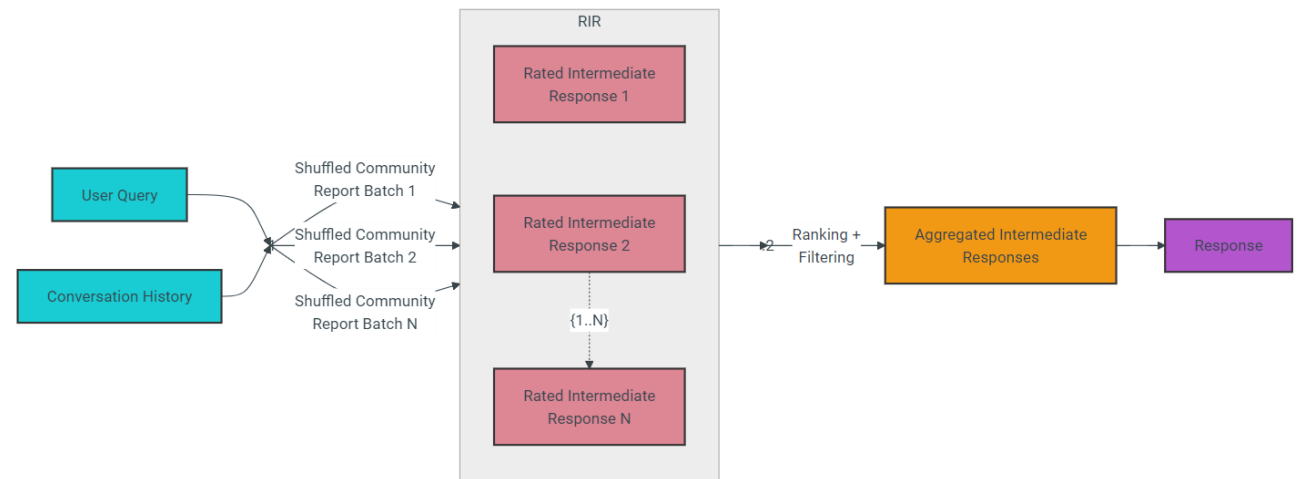
- **Entity-based Reasoning**
- **Local Search Method:** Augments LLM context by combining structured knowledge graph data with unstructured input documents.
- **Entity Extraction Process:**
  - **Identify Entities:** Extracts entities from the knowledge graph related to the query.
  - **Access Points:** Uses these entities to retrieve connected details, relationships, and relevant text from input documents.
- **Ideal for Entity-Specific Queries:** Suitable for questions about specific entities, such as "What are the healing properties of chamomile?"



Local Search Dataflow, Microsoft, [https://microsoft.github.io/graphrag/query/local\\_search/](https://microsoft.github.io/graphrag/query/local_search/)

# Global Search

- **Whole Dataset Reasoning**
- **Global Search Method:** Aggregates information from community reports in the knowledge graph.
- **Map-Reduce Process:**
  - **Map Step:** Breaks down reports into smaller chunks to generate intermediate responses.
  - **Reduce Step:** Combines key points to form a final summary.
- **Ideal for Overview Queries:** Useful for questions like "What are the top 5 themes?" or queries requiring a broad dataset overview.



Global Search Dataflow, Microsoft, [https://microsoft.github.io/graphrag/query/global\\_search/](https://microsoft.github.io/graphrag/query/global_search/)

# Hands-On 1

Adapting Graph RAG Artifacts to RDF Knowledge Graph

# Hands-On 2

Using SPARQL with a Knowledge Graph for RAG